

Training Graph Neural Networks with 1000 Layers

3/4/2025

Arun Cheriakara Joseph,
Faculty Of Math



Motivation

- Graphs Are Everywhere
- Real-World Examples
 - Netflix, Facebook, Amazon, Google Scholar
 - Internet = Vast Web Graph

Why Graph Neural Networks (GNNs)?

- Designed for graph-structured data
- Message passing between nodes
- Impressive results on **small** graphs

The Scaling Problem

- Memory bottleneck for large graphs
- Deep GNNs \rightarrow More parameters \rightarrow More memory
- Sampling/partitioning exist but not always ideal

Inspiration from NLP & Efficient Architectures

- NLP success: GPT, BERT, GPT-3 (scale improves performance)
- Efficient methods in NLP:
 - Reversible networks (RevNets)
 - Weight tying & equilibrium models

Proposed Methods

- Reversible Connections
- Weight-Tied Networks
- Graph Equilibrium (Implicit) GNN

Setup

- Input: A graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$
- $N=|V|$, $M=|E|$
- Adjacency matrix $A \in \mathbb{R}^{N \times N}$
- Node feature matrix $X \in \mathbb{R}^{N \times D}$
- (Optional) Edge feature matrix $U \in \mathbb{R}^{M \times F}$
- GNN operator: $f_w(X, A, U) \rightarrow X'$

Over-parameterized GNNs

- Deep GNNs often use **residual connections**
 - $X' = f_w(X, A, U) + X$
- Memory cost for L-layer GNN: $O(LND)$
 - L is the number of GNN layers,
 - N is the number of vertices,
 - D is the size of vertex features.
- Activation storage > parameter storage

Grouped Reversible GNN

- Inspired by **RevNets** + grouped convolutions
- Split features into C groups
- Reversible updates \rightarrow No storing intermediate states
- Memory complexity $\rightarrow O(ND)$

Forward Pass in Grouped Reversible GNNs

- Divide input vertex feature matrix X into C groups.
- Compute an exchange term
- Process each group iteratively

$$X'_0 = \sum_{i=2}^C X_i \quad (2)$$

$$X'_i = f_{w_i}(X'_{i-1}, A, U) + X_i, \quad i \in \{1, \dots, C\}, \quad (3)$$

Efficient Backpropagation

- Recompute X_i
- Recompute X'_0
- Reconstruct X_1

$$X_i = X'_i - f_{w_i}(X'_{i-1}, A, U), \quad i \in \{2, \dots, C\} \quad (4)$$

$$X'_0 = \sum_{i=2}^C X_i \quad (5)$$

$$X_1 = X'_1 - f_{w_1}(X'_0, A, U). \quad (6)$$

Avoiding Dropout Issues

- Dropout adds randomness, making reconstruction difficult.
- Solution: Use **shared dropout patterns** across all layers.
- This prevents memory from scaling with depth.

Memory Complexity: $O(ND)$

Memory Complexity & Practical Benefits

- Activations: $O(ND)$, not $O(LND)$
- Works with GCN, GAT, etc.
- Enables 1000+ layers on one GPU

Weight-Tied GNNs: The Idea

- Instead of learning different weights for each layer, we **reuse the same function** across layers.
- Reduces parameter count, making the model more **memory-efficient** and **generalizable**.

$$f_w^{(1)} := f_w^{(2)} \dots := f_w^{(L)}, \quad (9)$$

Two Types of Weight-Tied GNNs

- **Weight-Tied Residual GNNs:**
 1. Shares weights across layers.
 2. **Still requires storing activations**, so memory complexity remains $O(LND)$.
- **Weight-Tied Reversible GNNs:**
 1. Combines weight-tying with **reversibility**.
 2. Memory is reduced to $O(ND)$ (independent of depth).

Extending Weight-Tying to Groups

- Instead of weight-tying for all features, we divide features into **C groups**.
- Each group has its own tied function:

$$f_{w_i}^{(1)} := f_{w_i}^{(2)} \dots := f_{w_i}^{(L)}, i \in \{1, \dots, C\} \quad (10)$$

- Helps retain expressiveness while reducing parameters.

Weight-Tied Memory Complexity

Model Type	Memory Complexity
Regular Deep GNN	$O(LND)$
Weight-Tied Residual GNN	$O(LND)$
Weight-Tied Reversible GNN	$O(ND)$

Deep Equilibrium GNN: Overview

- Instead of explicitly stacking layers, the model **converges to a fixed point**.
- The equilibrium equation:

$$Z^* = f_w^{\text{DEQ}}(Z^*, X, A, U), \quad (11)$$

- The network **iterates until it reaches this stable state**.
- No need to define a fixed depth for the model.

DEQ-GNN Equations

- The forward pass of DEQ-GNN is implemented with a root-finding algorithm (e.g. Broyden's method)
- The gradients are obtained by implicitly differentiating through the equilibrium node state for the backward pass.

$$Z' = \text{GraphConv}(Z_{\text{in}}, A, U) \quad (12)$$

$$Z'' = \text{Norm}(Z' + X) \quad (13)$$

$$Z''' = \text{GraphConv}(\text{Dropout}(\text{ReLU}(Z'')), A, U) \quad (14)$$

$$Z_o = \text{Norm}(\text{ReLU}(Z''' + Z')), \quad (15)$$

Memory Efficiency in DEQ-GNNs

Model Type	Memory Complexity
Regular Deep GNN	$O(LND)$
DEQ-GNN	$O(ND)$

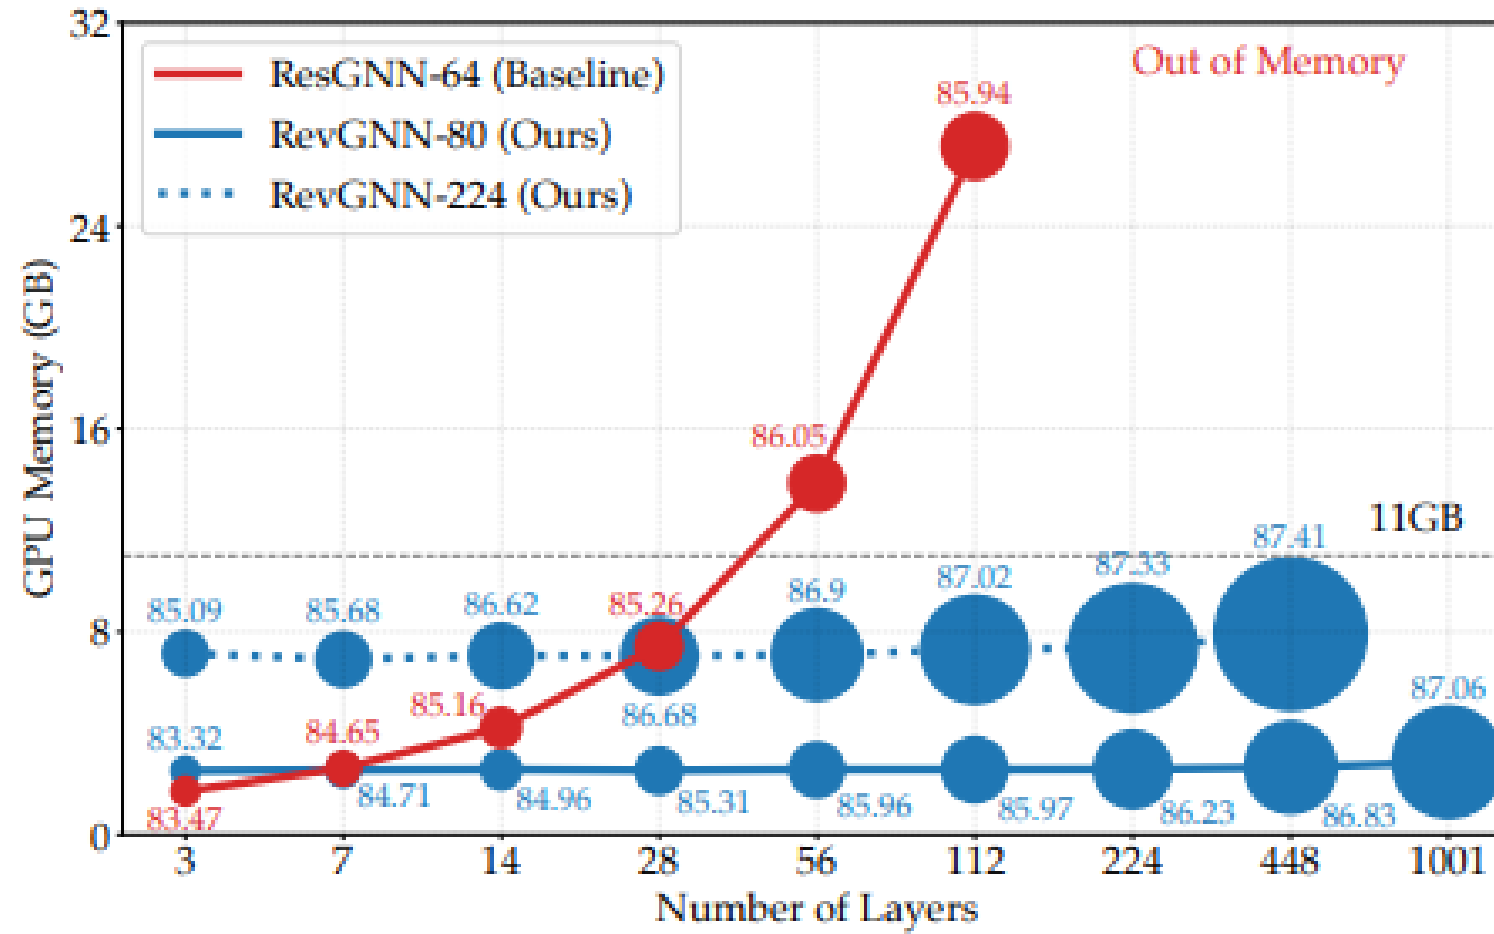
Analysis of Different Deep GNN Architectures

- **Dataset:** ogbn-proteins (OGB benchmark)
- **Training method:** Mini-batch training with random partitioning
- **Baseline Model:** ResGNN (Li et al., 2020)
- **Metric:** ROC-AUC score (higher is better)

Baseline GNN – ResGNN / Reversible GNN – RevGNN

- **Baseline ResGNN**
 - It's a **pre-activation residual** GNN
- **Reversible GNN - RevGNN**
 - No need to store all activations
 - Constant memory usage
 - Scales beyond 1000 layers


ResGNN VS RevGNN




WT-ResGNN vs. WT-RevGNN

- **Sharing the same parameters** across all layers
- Parameter count **does not grow** with depth
- **WT-ResGNN**: Residual, weight-tied
- **WT-RevGNN**: Reversible, weight-tied
- Same number of parameters, **different memory usage**

WT-ResGNN vs. WT-RevGNN



Equilibrium GNN



Over-parameterized Deep GNNs

- Experiments on **OGB** datasets (ogbn-proteins, ogbn-arxiv)
- Show **state-of-the-art** results with reversible GNN

ogbn-Protein Results

Table 1. Results on the ogbn-proteins dataset compared to SOTA. RevGNN-Deep has 1001 layers with 80 channels each. It achieves SOTA performance with minimal GPU memory for training. RevGNN-Wide has 448 layers with 224 channels each. It achieves the best accuracy while consuming a moderate amount of GPU memory.

Model	ROC-AUC \uparrow	Mem \downarrow	Params
GCN (Kipf & Welling)	72.51 ± 0.35	4.68	96.9k
GraphSAGE (Hamilton et al.)	77.68 ± 0.20	3.12	193k
DeeperGCN (Li et al.)	86.16 ± 0.16	27.1	2.37M
UniMP (Shi et al.)	86.42 ± 0.08	27.2	1.91M
GAT (Veličković et al.)	86.82 ± 0.21	6.74	2.48M
UniMP+CEF (Shi et al.)	86.91 ± 0.18	27.2	1.96M
Ours (RevGNN-Deep)	87.74 ± 0.13	2.86	20.03M
Ours (RevGNN-Wide)	88.24 ± 0.15	7.91	68.47M

ogbn-arxiv Results

Table 2. Results on the ogbn-arxiv dataset compared to SOTA. RevGCN-Deep has 28 layers with 128 channels each. It achieves SOTA performance with minimal GPU memory. RevGAT-Wide has 5 layers with 1068 channels each. RevGAT-SelfKD denotes the student models with 5 layers and 768 channels. It achieves the best accuracy while consuming a moderate amount of GPU memory.

Model	ACC \uparrow	Mem \downarrow	Params
GraphSAGE (Hamilton et al.)	71.49 ± 0.27	1.99	219k
GCN (Kipf & Welling)	71.74 ± 0.29	1.90	143k
DAGNN (Liu et al.)	72.09 ± 0.25	2.40	43.9k
DeeperGCN (Li et al.)	72.32 ± 0.27	21.6	491k
GCNII (Chen et al.)	72.74 ± 0.16	17.0	2.15M
GAT (Veličković et al.)	73.91 ± 0.12	5.52	1.44M
UniMP_v2 (Shi et al.)	73.97 ± 0.15	25.0	687k
Ours (RevGCN-Deep)	73.01 ± 0.31	1.84	262k
Ours (RevGAT-Wide)	74.05 ± 0.11	8.49	3.88M
Ours (RevGAT-SelfKD)	74.26 ± 0.17	6.60	2.10M

ogbn-arxiv Results on different GNN operators

Table 3. Results with different GNN operators on the *ogbn-arxiv*. All GAT models use label propagation. #L and #Ch denote the number of layers and channels respectively. *Baselines* are in italic.


Model	#L	#Ch	ACC \uparrow	Mem \downarrow	Params
<i>ResGCN</i>	28	128	72.46 ± 0.29	11.15	491k
RevGCN	28	128	73.01 ± 0.31	1.84	262k
RevGCN	28	180	73.22 ± 0.19	2.73	500k
<i>ResSAGE</i>	28	128	72.46 ± 0.29	8.93	950k
RevSAGE	28	128	72.69 ± 0.23	1.17	491k
RevSAGE	28	180	72.73 ± 0.10	1.57	953k
<i>ResGEN</i>	28	128	72.32 ± 0.27	21.63	491k
RevGEN	28	128	72.34 ± 0.18	4.08	262k
RevGEN	28	180	72.93 ± 0.10	5.67	500k
<i>ResGAT</i>	5	768	73.76 ± 0.13	9.96	3.87M
RevGAT	5	768	74.02 ± 0.18	6.30	2.10M
RevGAT	5	1068	74.05 ± 0.11	8.49	3.88M

Analysis of Complexities

Table 4. Comparison of complexities. L is the number of layers, D is the number of hidden channels, N is of the number of nodes, B is the batch size of nodes and R is the number of sampled neighbors of each node. K is the maximum Broyden iterations.

Method	Memory	Params	Time
Full-batch GNN	$\mathcal{O}(LND)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
GraphSAGE	$\mathcal{O}(R^L BD)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(R^L ND^2)$
VR-GCN	$\mathcal{O}(LND)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2 + R^L ND^2)$
FastGCN	$\mathcal{O}(LRBD)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(RLND^2)$
Cluster-GCN	$\mathcal{O}(LBD)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
GraphSAINT	$\mathcal{O}(LBD)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
Weight-tied GNN	$\mathcal{O}(LND)$	$\mathcal{O}(D^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
RevGNN	$\mathcal{O}(ND)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
WT-RevGNN	$\mathcal{O}(ND)$	$\mathcal{O}(D^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
DEQ-GNN	$\mathcal{O}(ND)$	$\mathcal{O}(D^2)$	$\mathcal{O}(K \ A\ _0 D + KND^2)$
RevGNN + Subgraph Sampling	$\mathcal{O}(BD)$	$\mathcal{O}(LD^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
WT-RevGNN + Subgraph Sampling	$\mathcal{O}(BD)$	$\mathcal{O}(D^2)$	$\mathcal{O}(L \ A\ _0 D + LND^2)$
DEQ-GNN + Subgraph Sampling	$\mathcal{O}(BD)$	$\mathcal{O}(D^2)$	$\mathcal{O}(K \ A\ _0 D + KND^2)$

Results



UNIVERSITY OF WATERLOO



Thank You!