

SIGN: Scalable Inception Graph Neural Networks

Fabrizio Frasca et al.

February 28, 2025

Twitter, Imperial College London

Presented by: George Giapitzakis

Graph Neural Networks: A Brief Overview

- Graph Neural Networks (GNNs) extend deep learning methods to graph-structured data.
- Applications: social networks, chemistry, physics, recommender systems, etc.
- Standard GNNs often face scalability challenges when dealing with large graphs.

Scaling Challenges in GNNs

- Many GNNs rely on neighbor sampling or multi-layer message passing.
- Deep architectures can suffer from over-smoothing and high computational cost.
- Sampling may introduce bias and inefficiency in both training and inference.

Motivation: How can we design a GNN that scales to web-scale graphs without these drawbacks?

Motivation for a Sampling-Free Approach

- **Scalability:** Web-scale graphs (e.g., billions of edges) require methods whose cost is independent of the graph structure during training/inference.
- **Expressiveness:** Retaining performance without sacrificing the ability to capture local graph structure.
- **Efficiency:** Avoiding iterative neighborhood sampling can greatly reduce both runtime and memory usage.

Inspiration from Inception Modules

- In CNNs, Inception modules combine filters of various sizes to capture multi-scale features.
- SIGN adapts this idea for graphs: using multiple precomputed diffusion operators as “filters” to capture different scales of neighborhood information.
- This approach bypasses the need for deep stacking of convolutional layers.

The SIGN Architecture

- SIGN stands for **S**calable **I**Nception **G**raph **N**eural network.
- Core idea: precompute multiple graph diffusion operators and combine their outputs.
- The model is defined as:

$$\mathbf{Z} = \sigma\left([\mathbf{X}\Theta_0, \mathbf{A}_1\mathbf{X}\Theta_1, \dots, \mathbf{A}_r\mathbf{X}\Theta_r]\right)$$

$$\mathbf{Y} = \xi(\mathbf{Z}\Omega)$$

- \mathbf{X} : Input node features; \mathbf{A}_i : Diffusion operators; Θ_i, Ω : Learnable parameters.

Precomputation of Diffusion Operators

- Diffusion operators $\mathbf{A}_1, \dots, \mathbf{A}_r$ can be based on:
- Normalized adjacency (GCN-style)
 - Personalized PageRank (PPR)
 - Triangle-induced (motif-based)
- These operators are computed **once** before training.




Figure 1: The SIGN architecture for r generic graph filtering operators. \parallel denotes the concatenation operation.

Relation to Inception Modules in CNNs

- Similar to using multiple kernel sizes in CNNs, SIGN uses different operators to capture various scales.
- Each operator acts like a convolution filter with a different receptive field.
- This multi-scale feature extraction is key to handling heterogeneous graph structures.

- **Preprocessing:** $O(r|E|d)$, where r is the number of operators, $|E|$ is the number of edges, and d the feature dimension.
- **Training/Inference:** Complexity becomes $O(rL_{ff}Nd^2)$ (similar to an MLP), independent of the graph structure.
- **Key advantage:** Avoids repeated neighborhood expansion and reduces redundant computations.

→ Datasets:

→ *Inductive*: Reddit, Flickr, Yelp, PPI.

→ *Transductive*: ogbn-products, ogbn-papers100M.

→ *Scalability*: Wikipedia links (with synthetic features).

→ **Baselines**: GCN, GraphSAGE, ClusterGCN, GraphSAINT, S-GCN.

→ **Metrics**: F1 scores, accuracy, and runtimes (preprocessing, training, and inference).

Key results: Accuracy

- SIGN achieves state-of-the-art results on ogbn-papers100M (over 110M nodes, 1.6B edges).
- Outperforms sampling-based methods in both accuracy and speed on several benchmarks.

	Reddit	Flickr	PPI	Yelp
<i>GCN</i> [34]	0.933±0.000	0.492±0.003	0.515±0.006	0.378±0.001
<i>FastGCN</i> [11]	0.924±0.001	0.504±0.001	0.513±0.032	0.265±0.053
<i>Stochastic-GCN</i> [12]	0.964±0.001	0.482±0.003	0.963±0.010	0.640±0.002
<i>AS-GCN</i> [30]	0.958±0.001	0.504±0.002	0.687±0.012	—
<i>GraphSAGE</i> [24]	0.953±0.001	0.501±0.013	0.637±0.006	0.634±0.006
<i>ClusterGCN</i> [13]	0.954±0.001	0.481±0.005	0.875±0.004	0.609±0.005
<i>GraphSAINT</i> [64]	0.966±0.001	0.511±0.001	0.981±0.004	0.653±0.003
<i>S-GCN</i> [59]	0.949±0.000	0.502±0.001	0.892±0.015	0.358±0.006
<i>SIGN</i> (<i>p, s, t</i>)	0.968±0.000 (4, 2, 0)	0.514±0.001 (4, 0, 1)	0.970±0.003 (2, 0, 1)	0.631±0.003 (2, 0, 1)

Figure 2: Micro-averaged $F1$ scores. For SIGN, only the best performing configuration is shown. The top three performance scores are highlighted as: **First**, **Second**, **Third**.

Ablation Study: Operator Combinations

- Evaluated different configurations of simple, PPR-based, and triangle-based operators.
- Results indicate that dataset-specific operator selection is crucial:

	Reddit	Flickr	PPI	Yelp
<i>SIGN</i> (2,0,0)	0.966±0.003	0.503±0.003	0.965±0.002	0.623±0.005
<i>SIGN</i> (2,0,1)	0.966±0.000	0.510±0.001	0.970±0.003	0.631±0.003
<i>SIGN</i> (2,2,0)	0.967±0.000	0.495±0.002	0.964±0.003	0.617±0.005
<i>SIGN</i> (4,0,0)	0.967±0.000	0.508±0.001	0.959±0.002	0.623±0.004
<i>SIGN</i> (4,0,1)	0.967±0.000	0.514±0.001	0.965±0.003	0.623±0.004
<i>SIGN</i> (4,2,0)	0.968±0.000	0.500±0.001	0.930±0.010	0.618±0.004
<i>SIGN</i> (4,2,1)	0.967±0.000	0.508±0.002	0.969±0.001	0.620±0.004

Figure 3: Impact of various operator combinations on inductive datasets. Best results are in bold.

Ablation Study: Operator Combinations

	<i>Training</i>	<i>Validation</i>	<i>Test</i>
<i>MLP</i>	84.03 \pm 0.93	75.54 \pm 0.14	61.06 \pm 0.08
<i>Node2Vec</i>	93.39 \pm 0.10	90.32 \pm 0.06	72.49 \pm 0.10
<i>S-GCN(L=5)</i>	92.54 \pm 0.09	91.38 \pm 0.07	74.87 \pm 0.25
<i>ClusterGCN</i>	93.75 \pm 0.13	92.12 \pm 0.09	78.97\pm0.33
<i>GraphSAINT</i>	92.71 \pm 0.14	91.62 \pm 0.08	79.08\pm0.24
<i>SIGN(3,0,0)</i>	96.21 \pm 0.31	92.99\pm0.05	76.52 \pm 0.14
<i>SIGN(3,0,1)</i>	96.46\pm0.29	92.93 \pm 0.04	75.73 \pm 0.20
<i>SIGN(3,3,0)</i>	96.87\pm0.23	93.02\pm0.04	77.13 \pm 0.10
<i>SIGN(5,0,0)</i>	95.99 \pm 0.69	92.98 \pm 0.18	76.83 \pm 0.39
<i>SIGN(5,3,0)</i>	96.92\pm0.46	93.10\pm0.08	77.60\pm0.13

Figure 4: Results on ogbn-papers100M. The top three performance scores are highlighted as: **First**, **Second**, **Third**.

Runtime Analysis

- SIGN's preprocessing takes longer than some sampling methods but is executed only once.
- Training and inference times are drastically reduced.
- Convergence is faster, leading to overall time savings on large graphs.

	ogbn-products			Wikipedia		
	<i>Preprocessing</i>	<i>Training</i>	<i>Inference</i>	<i>Preprocessing</i>	<i>Training</i>	<i>Inference</i>
<i>ClusterGCN</i>	36.93 ± 0.52	13.34 ± 0.16	93.00 ± 0.68	—	—	183.76 ± 3.01
<i>GraphSAINT</i>	52.06 ± 0.54	2.89 ± 0.05	94.76 ± 0.81	123.60 ± 1.60	135.73 ± 0.06	209.86 ± 4.73
<i>SIGN-2</i>	88.21 ± 1.33	1.04 ± 0.10	2.86 ± 0.10	192.88 ± 0.12	62.37 ± 0.17	13.40 ± 0.15
<i>SIGN-4</i>	160.16 ± 1.20	1.54 ± 0.04	3.79 ± 0.08	326.21 ± 1.14	93.84 ± 0.08	18.15 ± 0.05
<i>SIGN-6</i>	226.48 ± 1.43	2.05 ± 0.00	4.84 ± 0.08	459.24 ± 0.14	125.24 ± 0.03	22.94 ± 0.02
<i>SIGN-8</i>	297.92 ± 2.92	2.53 ± 0.04	5.88 ± 0.09	598.67 ± 0.82	154.73 ± 0.12	27.69 ± 0.11

Figure 5: Preprocessing, Training, and Inference times.

- **Scalability:** SIGN scales to web-scale graphs without relying on sampling.
- **Efficiency:** Precomputing diffusion operators shifts the heavy computation offline.
- **Performance:** Competitive or state-of-the-art accuracy across a range of benchmarks.
- **Simplicity vs. Expressiveness:** A single graph convolutional layer with multiple diffusion operators can outperform deeper, sampling-based architectures.

- **Depth vs. Width:** Results suggest that wide (multi-operator) shallow architectures are effective.
- **Extensions:**
 - Incorporating more expressive local operators (e.g., attention mechanisms).
 - Exploring higher-order network motifs and temporal motifs.
 - Further optimization of precomputation for distributed systems.

Thank you!
Any Questions?